

Recitation 4

Recitation Instructor: Shivam Verma

Email: shivamverma@nyu.edu

Ph: 718-362-7836

Office hours: WWH 605 (2.50 - 4.50 pm, Tuesdays)

Brief overview from last week (15 min)

- Big Oh notation - $f(x) = O(g(x)) \Leftrightarrow |f(x)| \leq M|g(x)|, x \geq x_0$. (measures the limiting behaviour of a function)
- Big Oh notation used in numerical analysis to measure # of mathematical operations (flops) for a numerical computation
- Eg. Matrix-vector, matrix-matrix multiplication
- How many operations does calculating a determinant take using the usual method?
- Ans: $O(n!)$
- LU factorization of matrices

Today's Agenda

- Matrix and vector norms
- Discussion of Homework questions
 - Overview of Gaussian elimination
 - Householder's matrices
 - Extracting submatrices/vectors from a matrix
- Cholesky decomposition - algorithm, order of computation

Practicing solving linear equations in MATLAB (20 min)

- The '****' operator in MATLAB is used for solving all kinds of linear equations. Try it!
- **Forward substitution in MATLAB** (example)
- Comparing **three different implementations** of forward substitution in MATLAB - looping over rows and columns (forward1), looping over rows only (forward2), looping over columns only (forward3).

Try the following code in MATLAB.

- Forward1 uses two for loops to do forward substitution.
- Forward2 uses one for loop (looping over rows)
- Forward3 uses one for loop (looping over columns)

Run the tester code below and observe the time taken for each of the forward substitution processes.

<pre>function [x] = forward1(L,b) %naive forward substitution by looping %over rows and columns n = length(b); x = zeros(n,1); x(1) = b(1)/L(1,1); for i=2:n s=0; for j=2:n s = s+L(i,j-1)*x(j-1); end x(i) = (b(i)-s)/L(i,i); end end</pre>	<pre>function [x] = forward2(L,b) %looping over rows n = length(b); x = zeros(n,1); x(1) = b(1)/L(1,1); for i=2:n x(i) = (b(i)-L(i,1:i-1)*x(1:i-1))/L(i,i); end end</pre>	<pre>function [x] = forward3(L,b) %looping over columns n = length(b); for j=1:n-1 b(j) = b(j)/L(j,j); b(j+1:n) = b(j+1:n) - b(j)*L(j+1:n,j); end b(n) = b(n)/L(n,n); x=b; end</pre>
<pre>%test forward1,forward2,forward3 N=10000; A = rand(N); %generates NxN random matrix b = rand(N,1); L = tril(A); %converts A to lower triangular matrix tic %starts measuring time taken c0 = L\b; toc %stops measuring time taken tic c1 = forward1(L,b); toc tic c2 = forward2(L,b); toc tic c3 = forward3(L,b); toc</pre>		

- Why do the **timings** of these three implementations of the **same algorithm** **differ** so much?
- Ans.
 - The MATLAB '\ ' operator is the most efficient and takes least time, as expected.

- The ‘two for loops’ implementation (forward1) should take around 7 seconds, while forward2 and forward3 take around 0.75 and 0.3 seconds respectively.
- Despite the fact that each method requires roughly the **same number of floating point operations**, timings differ significantly. This is mainly due to **memory access** and how **matrices are stored in memory**. **MATLAB stores matrices as one-dimensional arrays, column by column**. You can see that by accessing matrices with only one index, i.e., using $A(k)$ for an $n \times n$ matrix returns the entry $A(m, l)$, where $k = m + (l - 1)n$, $1 \leq l, 1 \leq m \leq n$. Numbers that are next to each other in memory can be read from memory much faster than numbers that are stored further away from each other.
- **Lesson: Columns and rows are NOT the same in MATLAB!**
- Finally, LU decomposition in MATLAB.
 - $[L \ U \ P] = \text{lu}(A)$
 - **help lu**

Norms and inequalities (15 min)


- What is a norm?
- An L_p norm is defined as $\|x\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$
- 2-norm or Euclidean norm: $\|x\|_2 = \sqrt{\sum_{i=1}^n x_i^2}$
- 1-norm: $\|x\|_1 = \sum_{i=1}^n |x_i|$
- ∞ - norm: $\|x\|_\infty = \max(|x_1|, |x_2|, \dots, |x_n|)$

- Visualizing norms in MATLAB

```
function [] = VisualizeNorms(x_c,r,p)
%plots a p-norm ball with radius r, centred at x_c
    low = min(x_c);
    high = max(x_c);
    x = linspace(low-2*r,high+2*r);
    y = linspace(low-2*r,high+2*r);
    theta = 0:pi/100:3*pi;
    xp = r*sign(cos(theta)).*abs(cos(theta)).^(2/p);
    yp = r*sign(sin(theta)).*abs(sin(theta)).^(2/p);
    plot(xp,yp,'-', 'LineWidth',1.5)
end
```

```
%tester code
x_c = [0,0]; %center at origin
r = 1; %radius of ball is 1
plot(x_c(1),x_c(2),'k+') %plots the center
hold on
xlabel('x_1')
ylabel('x_2')
axis([-x_c(1)+2*r x_c(1)+2*r -(x_c(2)+2*r) x_c(2)+2*r])
axis equal
VisualizeNorms(x_c,r,0.5) %p = 0.5
VisualizeNorms(x_c,r,1) %p = 1
VisualizeNorms(x_c,r,1.5) %p = 1.5
VisualizeNorms(x_c,r,2) %p = 2
VisualizeNorms(x_c,r,inf) %p = inf
legend('x_c','p = 0.5','p = 1','p = 1.5','p = 2','p = \infty')
```

You should get a plot which looks like this:



- Think about why this happens. Read the references to know more about norms.
- Can also define matrix norms

$$\circ \sup_{v \in \mathbb{R}^n} \frac{\|Av\|}{\|v\|}$$

Special matrices

- Banded matrices
 - $a_{ij} = 0$ for all i and j such that $|i - j| < k$ for a k -band matrix
 - Eg. tri-diagonal matrix
- Positive-definite matrices
 - $x^T A x > 0$ for all non-zero x in real space
- What is a Householder matrix? Where is it used (hint: QR algorithm)? Why is it useful?

Discussion about Cholesky decomposition (15 min)

- What is a Cholesky decomposition? When can it be done? Why is it useful?
- If a matrix is positive semi-definite, it has a Cholesky decomposition $A = LL^T$.
- $Ax = b \Rightarrow LL^T x = b \Rightarrow L^T x = y$ (backward substitution) and $Ly = b$ (forward substitution)

write $A = A^{(1)}$ as

$$A^{(1)} = \left[\begin{array}{c|c} a_{11} & z^T \\ \hline z & B^{(1)} \end{array} \right],$$

where $z = (a_{12}, \dots, a_{1n})^T$ and after one elimination step we obtain

$$A^{(2)} = L_1 A^{(1)} = \left[\begin{array}{c|c} a_{11} & z^T \\ \hline 0 & B^{(2)} \end{array} \right] \quad \text{with} \quad L_1 = \left[\begin{array}{cccc} 1 & & & \\ -l_{21} & 1 & & \\ \vdots & & \ddots & \\ -l_{n1} & & & 1 \end{array} \right].$$

Now if we premultiply $A^{(2)}$ with L_1^T , then z^T in the first row is also eliminated and the remainder matrix $B^{(2)}$ remains unchanged, i.e.,

$$L_1 A^{(1)} L_1^T = \left[\begin{array}{c|ccc} a_{11} & 0 & \cdots & 0 \\ \hline 0 & & & \\ \vdots & & B^{(2)} & \\ 0 & & & \end{array} \right].$$

[Image source: Numerical Analysis in Modern Scientific Computing, Deuffhard & Hohmann (Ch. 1)]

- Cholesky algorithm -

```

for  $k := 1$  to  $n$  do
     $l_{kk} := (a_{kk} - \sum_{j=1}^{k-1} l_{kj}^2)^{1/2};$ 
    for  $i := k + 1$  to  $n$  do
         $l_{ik} = (a_{ik} - \sum_{j=1}^{k-1} l_{ij} l_{kj}) / l_{kk};$ 
    end for
end for

```

[Image source: Numerical Analysis in Modern Scientific Computing, Deuffhard & Hohmann (Ch. 1)]

- Computational cost = $n^3/3$ = half of LU decomposition

[Helpful resources](#)

1. [Why columns and rows are not the same in MATLAB](#)
2. [Notes on vector norms](#)
3. [Notes on matrix norms](#)
4. [Intuition behind norms](#)
5. [Cholesky decomposition explained](#)
6. [About householder transformation - 1](#)
7. [About householder transformation - 2](#)
8. *Numerical Mathematics (Quarteroni et al)* is a very good resource for LU/Cholesky/etc. It's also freely available for NYU students via Springer!